3.5. LABS



Exercise 3.4: Deploy A Simple Application

We will test to see if we can deploy a simple application, in this case the **nginx** web server.

1. Create a new deployment, which is a Kubernetes object, which will deploy an application in a container. Verify it is running and the desired number of containers matches the available.

```
student@cp:~$ kubectl create deployment nginx --image=nginx

deployment.apps/nginx created

student@cp:~$ kubectl get deployments

NAME READY UP-TO-DATE AVAILABLE AGE
nginx 1/1 1 1 8s
```

2. View the details of the deployment. Remember auto-completion will work for sub-commands and resources as well.

student@cp:~\$ kubectl describe deployment nginx

```
Name:
                       nginx
Namespace:
                       default
CreationTimestamp:
                       Mon, 23 Apr 2019 22:38:32 +0000
Labels:
                       app=nginx
Annotations:
                       deployment.kubernetes.io/revision: 1
Selector:
                       app=nginx
Replicas:
                       1 desired | 1 updated | 1 total | 1 ava....
StrategyType:
                       RollingUpdate
MinReadySeconds:
RollingUpdateStrategy: 25% max unavailable, 25% max surge
<output_omitted>
```

3. View the basic steps the cluster took in order to pull and deploy the new application. You should see several lines of output. The first column shows the age of each message, note that due to JSON lack of order the time LAST SEEN time does not print out chronologically. Eventually older messages will be removed.

```
student@cp:~$ kubectl get events

<output_omitted>
```

4. You can also view the output in **yaml** format, which could be used to create this deployment again or new deployments. Get the information but change the output to yaml. Note that halfway down there is status information of the current deployment.

```
student@cp:~$ kubectl get deployment nginx -o yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
annotations:
deployment.kubernetes.io/revision: "1"
creationTimestamp: 2017-09-27T18:21:25Z
coutput_omitted>
```



5. Run the command again and redirect the output to a file. Then edit the file. Remove the creationTimestamp, resourceVersion, and uid lines. Also remove all the lines including and after status:, which should be somewhere around line 120, if others have already been removed.

```
student@cp:~$ kubectl get deployment nginx -o yaml > first.yaml
student@cp:~$ vim first.yaml

<Remove the lines mentioned above>
```

6. Delete the existing deployment.

```
student@cp:~$ kubectl delete deployment nginx

deployment.apps "nginx" deleted
```

7. Create the deployment again this time using the file.

```
student@cp:~$ kubectl create -f first.yaml

deployment.apps/nginx created
```

8. Look at the yaml output of this iteration and compare it against the first. The creation time stamp, resource version and unique ID we had deleted are in the new file. These are generated for each resource we create, so we may need to delete them from yaml files to avoid conflicts or false information. You may notice some time stamp differences as well. The status should not be hard-coded either.

```
student@cp:~$ kubectl get deployment nginx -o yaml > second.yaml
student@cp:~$ diff first.yaml second.yaml
<output_omitted>
```

9. Now that we have worked with the raw output we will explore two other ways of generating useful YAML or JSON. Use the --dry-run option and verify no object was created. Only the prior nginx deployment should be found. The output lacks the unique information we removed before, but does have the same essential values.

```
student@cp:~$ kubectl create deployment two --image=nginx --dry-run=client -o yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
creationTimestamp: null
labels:
app: two
name: two
spec:
coutput_omitted>
```

student@cp:~\$ kubectl get deployment

```
NAME READY UP-TO-DATE AVAILABLE AGE
nginx 1/1 1 1 7m
```

Existing objects can be viewed in a ready to use YAML output. Take a look at the existing nginx deployment.

```
student@cp:~$ kubectl get deployments nginx -o yaml
```



3.5. LABS

```
apiVersion: apps/v1
kind: Deployment
metadata:
annotations:
deployment.kubernetes.io/revision: "1"
creationTimestamp: null
generation: 1
labels:
run: nginx

coutput_omitted>
```

11. The output can also be viewed in JSON output.

```
student@cp:~$ kubectl get deployment nginx -o json
```

12. The newly deployed **nginx** container is a light weight web server. We will need to create a service to view the default welcome page. Begin by looking at the help output. Note that there are several examples given, about halfway through the output.

```
student@cp:~$ kubectl expose -h

<output_omitted>
```

13. Now try to gain access to the web server. As we have not declared a port to use you will receive an error.

```
student@cp:~$ kubectl expose deployment/nginx
```

```
error: couldn't find port via --port flag or introspection
See 'kubectl expose -h' for help and examples.
```

14. To change an object configuration one can use subcommands apply, edit or patch for non-disruptive updates. The apply command does a three-way diff of previous, current, and supplied input to determine modifications to make. Fields not mentioned are unaffected. The edit function performs a get, opens an editor, then an apply. You can update API objects in place with JSON patch and merge patch or strategic merge patch functionality.

If the configuration has resource fields which cannot be updated once initialized then a disruptive update could be done using the replace --force option. This deletes first then re-creates a resource.

Edit the file. Find the container name, somewhere around line 31 and add the port information as shown below.

```
student@cp:~$ vim first.yaml
```





3

```
- image: nginx
imagePullPolicy: Always
name: nginx
ports: # Add these

- containerPort: 80 # three
protocol: TCP # lines

resources: {}
```

15. Due to how the object was created we will need to use replace to terminate and create a new deployment.

```
student@cp:~$ kubectl replace -f first.yaml

deployment.apps/nginx replaced
```

16. View the Pod and Deployment. Note the AGE shows the Pod was re-created.

```
student@cp:~$ kubectl get deploy,pod
```

```
NAME
                              READY
                                      UP-TO-DATE
                                                   AVAILABLE
                                                               AGE
deployment.apps/nginx
                        1/1
                               1
                                             1
                                                         2m4s
                                     STATUS
NAME
                             READY
                                               RESTARTS
                                                          AGE
pod/nginx-7db75b8b78-qjffm
                             1/1
                                     Running
```

17. Try to expose the resource again. This time it should work.

```
student@cp:~$ kubectl expose deployment/nginx
```

```
service/nginx exposed
```

18. Verify the service configuration. First look at the service, then the endpoint information. Note the ClusterIP is not the current endpoint. Calico provides the ClusterIP. The Endpoint is provided by kubelet and kube-proxy. Take note of the current endpoint IP. In the example below it is 192.168.1.5:80. We will use this information in a few steps.

```
student@cp:~$ kubectl get svc nginx
```

```
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
nginx ClusterIP 10.100.61.122 <none> 80/TCP 3m
```

student@cp:~\$ kubectl get ep nginx

```
NAME ENDPOINTS AGE
nginx 192.168.1.5:80 26s
```

19. Determine which node the container is running on. Log into that node and use **tcpdump**, which you may need to install using **apt-get install**, to view traffic on the tunl0, as in tunnel zero, interface. The second node in this example. You may also see traffic on an interface which starts with cali and some string. Leave that command running while you run **curl** in the following step. You should see several messages go back and forth, including a HTTP HTTP/1.1 200 OK: and a ack response to the same sequence.

```
Node: worker/10.128.0.5
```



3.5. LABS 5

```
student@worker:~$ sudo tcpdump -i tunl0
```

```
tcpdump: verbose output suppressed, use -v or -vv for full protocol...
listening on tunl0, link-type EN10MB (Ethernet), capture size...
<output_omitted>
```

20. Test access to the Cluster IP, port 80. You should see the generic nginx installed and working page. The output should be the same when you look at the ENDPOINTS IP address. If the **curl** command times out the pod may be running on the other node. Run the same command on that node and it should work.

```
student@cp:~$ curl 10.100.61.122:80
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
<output_omitted>
```

```
student@cp:~$ curl 192.168.1.5:80
```

21. Now scale up the deployment from one to three web servers.

student@cp:~\$ kubectl get deployment nginx

```
NAME READY UP-TO-DATE AVAILABLE AGE
nginx 1/1 1 1 12m
```

student@cp:~\$ kubectl scale deployment nginx --replicas=3

```
deployment.apps/nginx scaled
```

student@cp:~\$ kubectl get deployment nginx

```
NAME READY UP-TO-DATE AVAILABLE AGE
nginx 3/3 3 3 12m
```

22. View the current endpoints. There now should be three. If the UP-TO-DATE above said three, but AVAILABLE said two wait a few seconds and try again, it could be slow to fully deploy.

student@cp:~\$ kubectl get ep nginx

```
NAME ENDPOINTS AGE
nginx 192.168.0.3:80,192.168.1.5:80,192.168.1.6:80 7m40s
```

23. Find the oldest pod of the **nginx** deployment and delete it. The Tab key can be helpful for the long names. Use the AGE field to determine which was running the longest. You may notice activity in the other terminal where **tcpdump** is running, when you delete the pod. The pods with 192.168.0 addresses are probably on the cp and the 192.168.1 addresses are probably on the worker

student@cp:~\$ kubectl get pod -o wide

```
READY
                                 STATUS
                                           RESTARTS
NAME.
                                                     AGF.
                       1/1
                                                     14m 192.168.1.5
nginx-1423793266-7f1qw
                                 Running
                                          0
                                                     86s 192.168.1.6
nginx-1423793266-8w2nk
                       1/1
                                 Running
                                          0
                                                     86s 192.168.0.3
nginx-1423793266-fbt4b
                       1/1
                                          0
                                 Running
```

student@cp:~\$ kubectl delete pod nginx-1423793266-7f1qw



```
pod "nginx-1423793266-7f1qw" deleted
```

24. Wait a minute or two then view the pods again. One should be newer than the others. In the following example nine seconds instead of four minutes. If your **tcpdump** was using the veth interface of that container it will error out. Also note we are using a short name for the object.

student@cp:~\$ kubectl get po

```
NAME
                         READY
                                   STATUS
                                             RESTARTS
                                                        AGE
nginx-1423793266-13p69
                         1/1
                                   Running
                                            0
                                                        9s
nginx-1423793266-8w2nk
                         1/1
                                   Running
                                            0
                                                        4m1s
nginx-1423793266-fbt4b
                         1/1
                                   Running
                                            0
                                                        4m1s
```

25. View the endpoints again. The original endpoint IP is no longer in use. You can delete any of the pods and the service will forward traffic to the existing backend pods.

```
student@cp:~$ kubectl get ep nginx
```

```
NAME ENDPOINTS AGE
nginx 192.168.0.3:80,192.168.1.6:80,192.168.1.7:80 12m
```

26. Test access to the web server again, using the ClusterIP address, then any of the endpoint IP addresses. Even though the endpoints have changed you still have access to the web server. This access is only from within the cluster. When done use **ctrl-c** to stop the **tcpdump** command.

```
student@cp:~$ curl 10.100.61.122:80
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
    body
<output_omitted>
```

